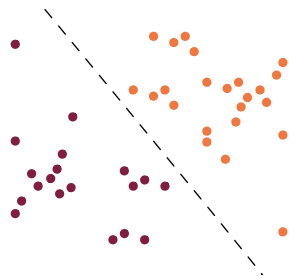
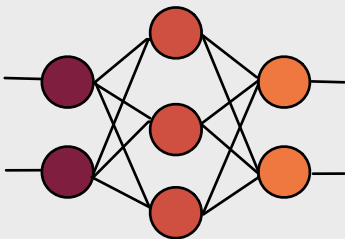


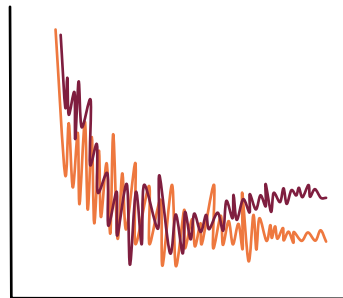
Overview



1. Find some data



2. Design a network



3. Train the network

Neural Architecture Search (NAS)

Algorithm 1 Standard NAS

generator = RNN()

for $i=1 : N$ **do**

net = generator.generate()

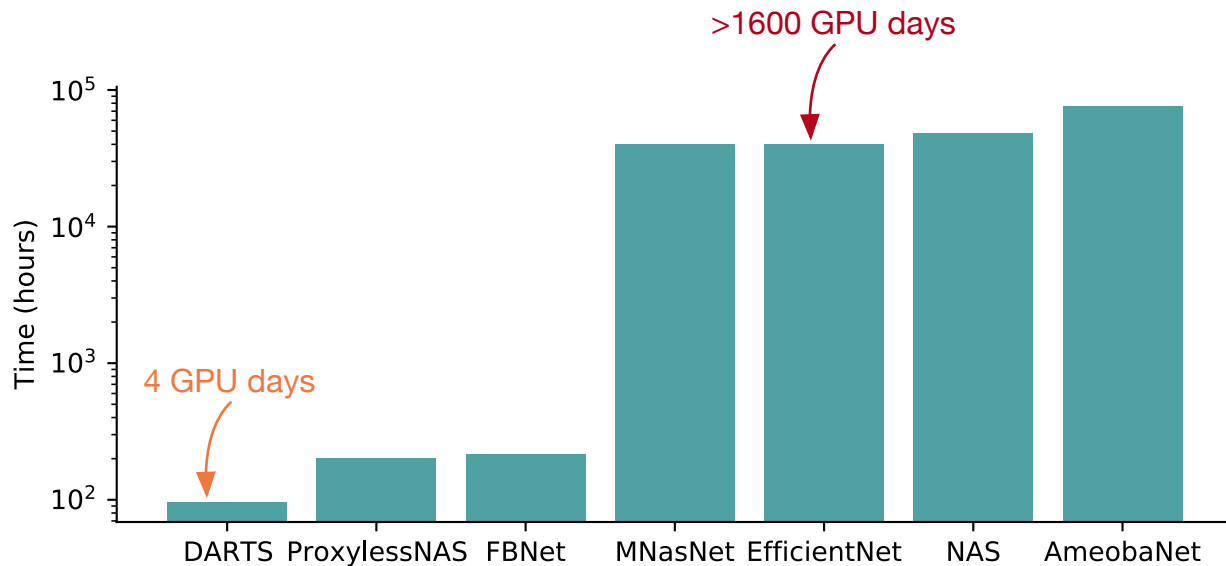
trained_net = net.train()

▷ *Training a net every step is expensive*

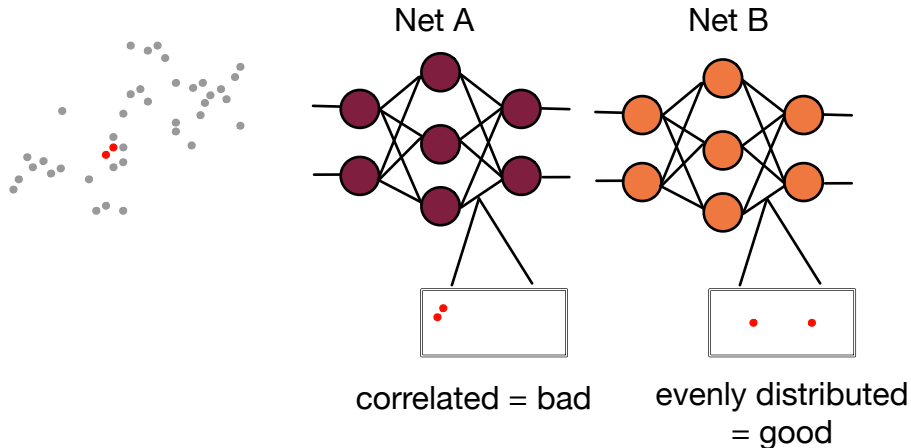
generator.update(trained_net)

chosen_net = generator.generate()

NAS is slow




Is there a property of networks *without training* that we can use?

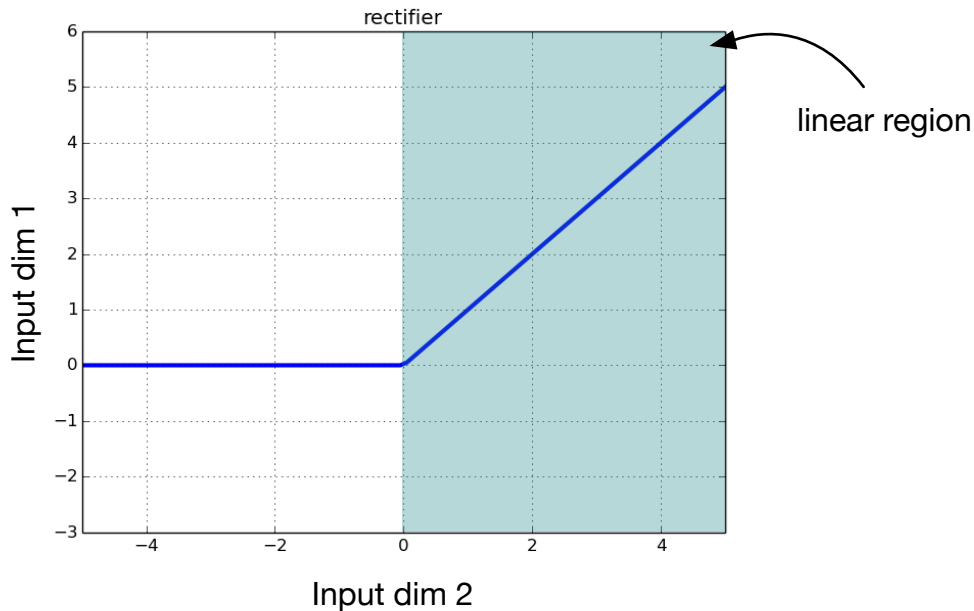
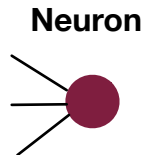


Linear Regions in Neural Networks

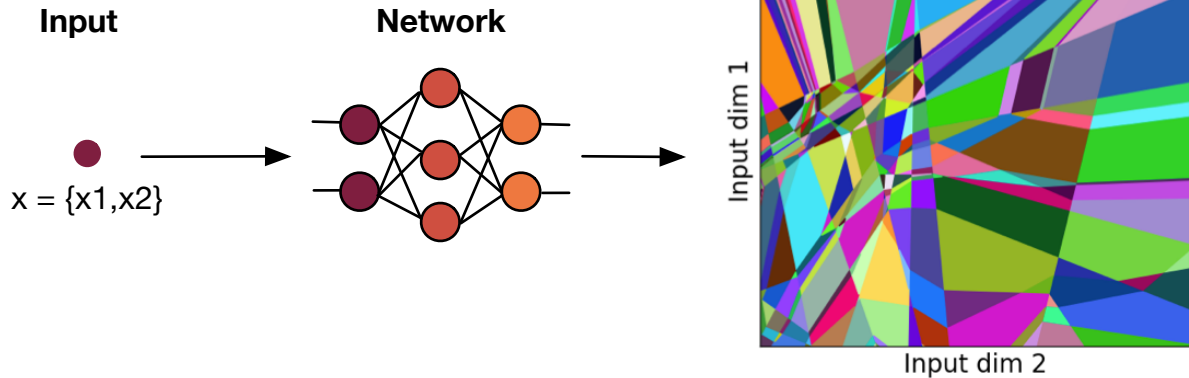
Input



$x = \{x_1, x_2\}$

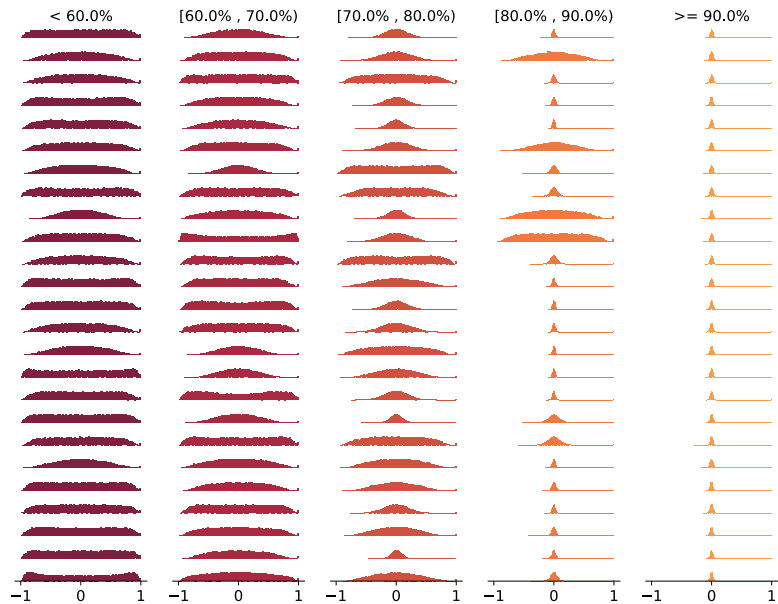


Linear Regions

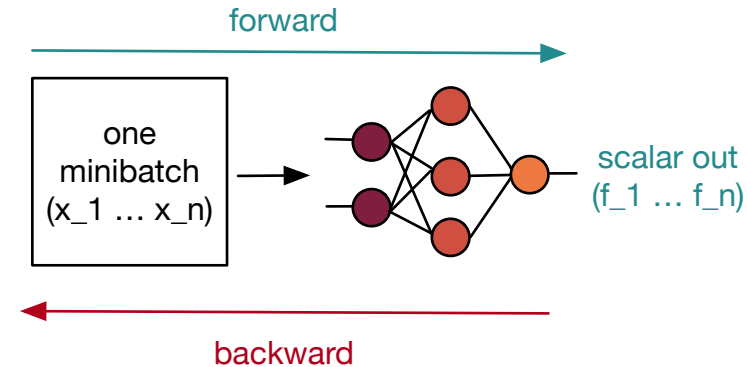


[1] Hanin, B. and Rolnick, D., 2019. Deep relu networks have surprisingly few activation patterns. In *Advances in Neural Information Processing Systems* (pp. 361-370).

Correlating Linear Regions



Correlating Linear Regions



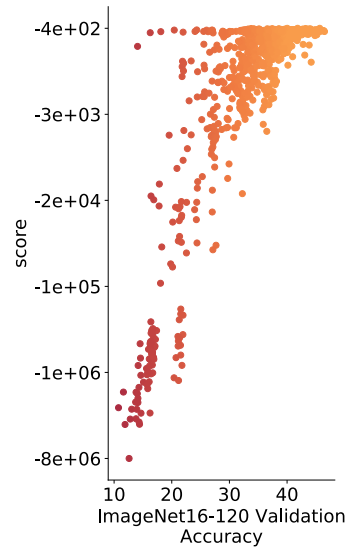
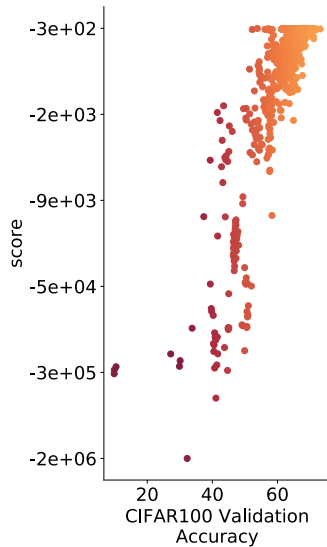
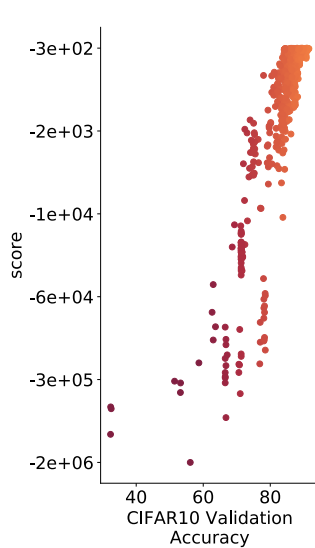
$$\mathbf{J} = \begin{pmatrix} \frac{\partial f_1}{\partial \mathbf{x}_1} & \frac{\partial f_2}{\partial \mathbf{x}_2} & \dots & \frac{\partial f_N}{\partial \mathbf{x}_N} \end{pmatrix}^\top$$

Scoring Jacobians

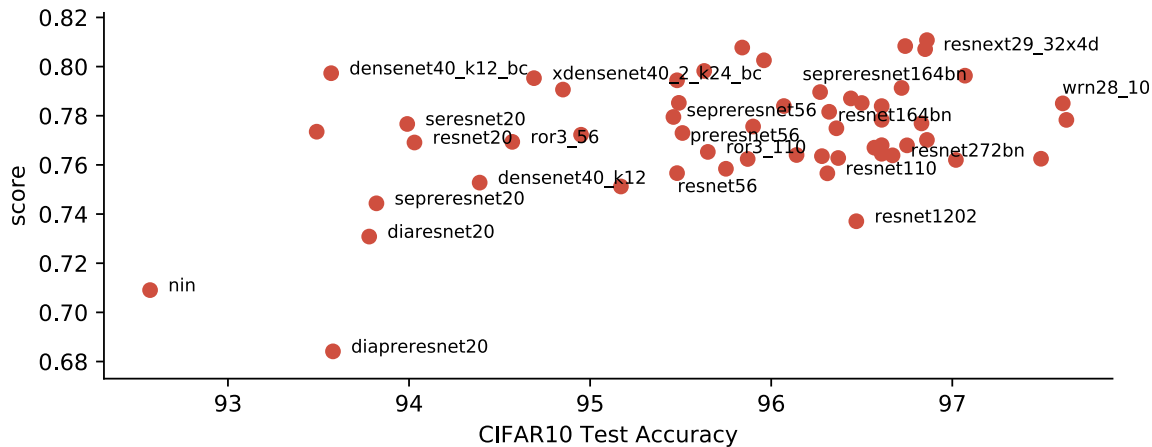
1. Get the input-output Jacobian (J) for one minibatch of data
2. Take eigenvalues of the correlation matrix of J
3. Score by KL divergence between two Gaussians:
 - one with 2. as the kernel
 - one with an uncorrelated kernel

```
def eval_score(jacob):  
    corrs = np.corrcoef(jacob)  
    v, _ = np.linalg.eig(corrs)  
    k = 1e-5  
    return -np.sum(np.log(v + k) + 1./(v + k))
```

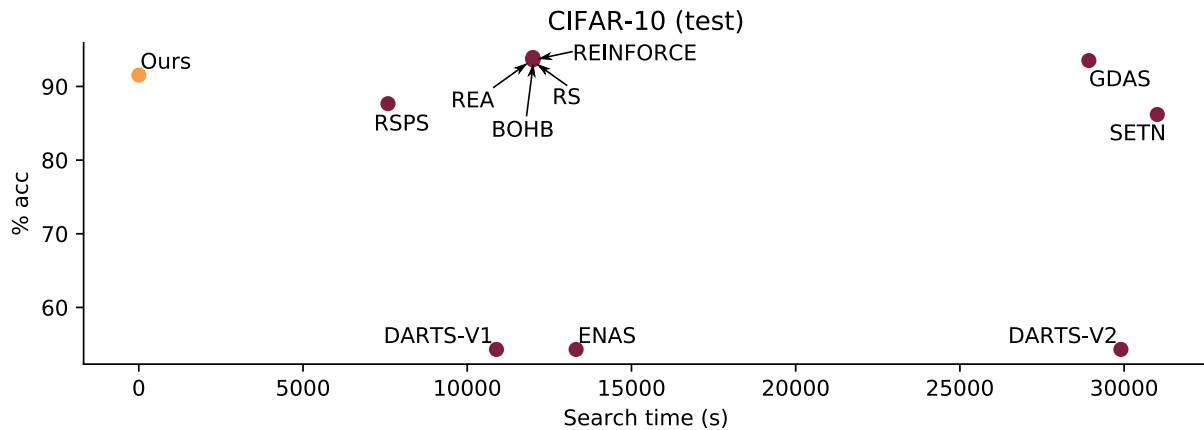
Results



Results



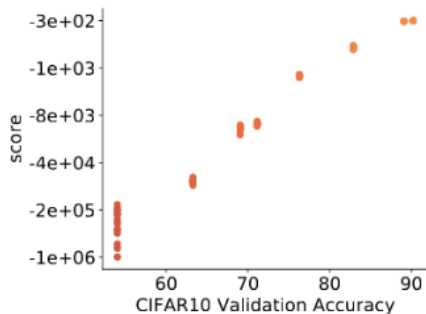
Results



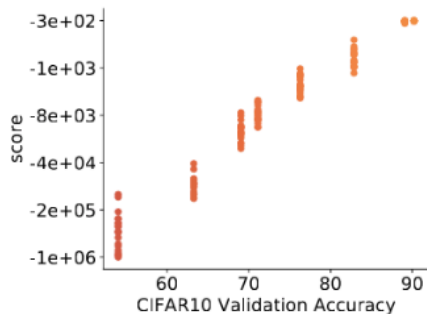
Questions?

Code: <https://github.com/BayesWatch/nas-without-training>

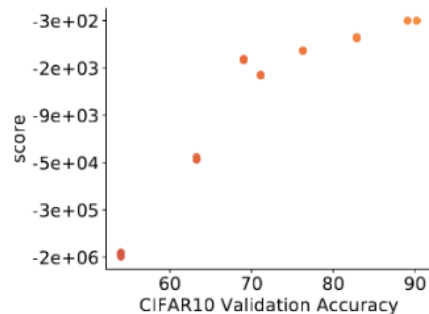
Control studies



(a) vary minibatch



(b) vary initialisation



(c) vary white noise input